

```
from google.colab import drive
```

```
drive.mount("/content/drive/")
```

```
↳ Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount(force=True)
```

```
pip install pyforest
```

```
↳ Requirement already satisfied: pyforest in /usr/local/lib/python3.6/dist-packages (1.0.0)
```

```
import pyforest
```

```
lazy_imports()
```

```
↳ ['import os',  
    'from sklearn.ensemble import RandomForestClassifier',  
    'import nltk',  
    'from pathlib import Path',  
    'from sklearn.ensemble import RandomForestRegressor',  
    'import pandas as pd',  
    'from openpyxl import load_workbook',  
    'import matplotlib as mpl',  
    'from sklearn.preprocessing import OneHotEncoder',  
    'import keras',  
    'import tqdm',  
    'import pydot',  
    'import spacy',  
    'import sys',  
    'from sklearn.feature_extraction.text import TfidfVectorizer',  
    'from sklearn.ensemble import GradientBoostingClassifier',  
    'import statistics',  
    'import tensorflow as tf',  
    'import plotly.express as px',  
    'import plotly as py',  
    'import pickle',  
    'import re',  
    'import glob',  
    'from pyspark import SparkContext',  
    'import numpy as np',  
    'import dash',  
    'import sklearn',  
    'from sklearn import svm',  
    'from sklearn.ensemble import GradientBoostingRegressor',  
    'from dask import dataframe as dd',  
    'from sklearn.manifold import TSNE',  
    'import datetime as dt',  
    'import seaborn as sns',  
    'import gensim',  
    'import matplotlib.pyplot as plt',  
    'import bokeh',  
    'import altair as alt',  
    'from sklearn.model_selection import train_test_split',  
    'import plotly.graph_objs as go']
```

```
import sqlite3
```

```
con=sqlite3.connect("/content/drive/My Drive/colab folder/Datasets/amazon-fine-food-reviews.db")
```

```
database=pd.read_sql_query("select * from reviews where score<=3",con)
```



```
database.Score.value_counts()
```



```
5    363122
4    80655
1    52268
2    29769
Name: Score, dtype: int64
```

```
def partition(x):
    if x>3:
        return 1
    return 0
```

```
sorted_database=database.Score.map(partition)
```

```
database.Score=sorted_database
```

```
database.Score.value_counts()
```



```
1    443777
0    82037
Name: Score, dtype: int64
```

```
sort=database.sort_values("ProductId")
```

```
sort
```



| | Id | ProductId | UserId | ProfileName | HelpfulnessNumerator | HelpfulnessDenominator |
|---------------|-----------|------------------|----------------|-----------------------------|-----------------------------|-------------------------------|
| 138706 | 150524 | 0006641040 | ACITT7DI6IDDL | shari zychinski | 0 | 0 |
| 138688 | 150506 | 0006641040 | A2IW4PEEKO2R0U | Tracy | 1 | 1 |
| 138689 | 150507 | 0006641040 | A1S4A3IQ2MU7V4 | sally sue "sally sue" | 1 | 1 |
| 138690 | 150508 | 0006641040 | AZGXZ2UUK6X | Catherine Hallberg "(Kate)" | 1 | 1 |
| 138691 | 150509 | 0006641040 | A3CMRKGE0P909G | Teresa | 3 | 3 |
| ... | ... | ... | ... | ... | ... | ... |
| 176791 | 191721 | B009UOFTUI | AJVB004EB0MVK | D. Christofferson | 0 | 0 |
| 1362 | 1478 | B009UOFU20 | AJVB004EB0MVK | D. Christofferson | 0 | 0 |
| 303285 | 328482 | B009UUS05I | ARL20DSHGVM1Y | Jamie | 0 | 0 |

```
final=sort.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=True)
final.shape
```

🔗 (364173, 10)

```
final.shape
```

```
↳ (261172, 10)
```

```
sort.shape
```

```
↳ (525814, 10)
```

```
final.Score.value_counts()
```

```
↳ 1    307063  
   0     57110  
   Name: Score, dtype: int64
```

▼ PREPROCESSING OR TEXT SUMMARIZATION

```
from bs4 import BeautifulSoup
```

```
import re
```

```
def decontracted(phrase):  
    # specific  
    phrase = re.sub(r"won't", "will not", phrase)  
    phrase = re.sub(r"can't", "can not", phrase)
```

```
    # general  
    phrase = re.sub(r"n't", " not", phrase)  
    phrase = re.sub(r"\ 're", " are", phrase)  
    phrase = re.sub(r"\ 's", " is", phrase)  
    phrase = re.sub(r"\ 'd", " would", phrase)  
    phrase = re.sub(r"\ 'll", " will", phrase)  
    phrase = re.sub(r"\ 't", " not", phrase)  
    phrase = re.sub(r"\ 've", " have", phrase)  
    phrase = re.sub(r"\ 'm", " am", phrase)  
    return phrase
```

```
stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you',  
    "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'her',  
    "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'theirs',  
    'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'am', 'is',  
    'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'doing', 'a',  
    'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'at', 'by', 'for',  
    'with', 'about', 'against', 'between', 'into', 'through', 'during', 'above', 'below',  
    'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'then', 'once', 'here',  
    'there', 'when', 'where', 'why', 'how', 'all', 'any', 'each', 'both', 'most', 'other',  
    'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very', 's', 't', 'can',  
    'will', 'just', 'don', "don't", 'should', "should've", 'now', 've', 'y', 'ain',  
    'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't",  
    'hadn', "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma',  
    'mightn', "mightn't", 'mustn', "mustn't", 'shan', "shan't", 'shouldn', "shouldn't",  
    'wasn', "wasn't", 'wouldn', "wouldn't"])
```

```
%%time  
from tqdm import tqdm  
preprocessed_reviews = []  
# tqdm is for printing the status bar  
for sentence in tqdm(final['Text'].values):  
    sentence = re.sub(r"http\S+", "", sentence)
```

```

sentence = BeautifulSoup(sentence, 'lxml').get_text()
sentence = decontracted(sentence)
sentence = re.sub("\S*\d\S*", "", sentence).strip()
sentence = re.sub('[^A-Za-z]+', ' ', sentence)
# https://gist.github.com/sebleier/554280
sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
preprocessed_reviews.append(sentence.strip())

```

```

↳ 100%|██████████| 364173/364173 [02:13<00:00, 2731.14it/s]CPU times: user 2min 8s, sys:
Wall time: 2min 13s

```

```

string = " geeks for geeks "

```

```

# prints the string without stripping
print(string)

```

```

# prints the string by removing leading and trailing whitespaces
print(string.strip())

```

```

↳ geeks for geeks
   geeks for geeks

```

```

# prints the string by removing geeks
print(string.strip(' geeks'))

```

```

↳ for

```

```

from tqdm import tqdm
preprocessed_summary = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Summary'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_summary.append(sentence.strip())

```

```

↳ 6%|███████| 22339/364173 [00:05<01:17, 4436.65it/s]/usr/local/lib/python3.6/dist-p
' Beautiful Soup.' % markup)
100%|██████████| 364173/364173 [01:22<00:00, 4439.45it/s]

```

```

final["Cleaned_text"]=preprocessed_reviews
final["Cleaned_summary"]=preprocessed_summary

```

```

↳

```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/10min.html  
    """Entry point for launching an IPython kernel.  
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/10min.html
```

final



| | Id | ProductId | UserId | ProfileName | HelpfulnessNumerator | HelpfulnessDenominator |
|--------|--------|------------|----------------|-----------------------------|----------------------|------------------------|
| 138706 | 150524 | 0006641040 | ACITT7DI6IDDL | shari zychinski | 0 | 0 |
| 138688 | 150506 | 0006641040 | A2IW4PEEKO2R0U | Tracy | 1 | 1 |
| 138689 | 150507 | 0006641040 | A1S4A3IQ2MU7V4 | sally sue "sally sue" | 1 | 1 |
| 138690 | 150508 | 0006641040 | AZGXZ2UUK6X | Catherine Hallberg "(Kate)" | 1 | 1 |
| 138691 | 150509 | 0006641040 | A3CMRKGE0P909G | Teresa | 3 | 3 |
| ... | ... | ... | ... | ... | ... | ... |
| 178145 | 193174 | B009RSR8HO | A4P6AN2L435PV | romarc | 0 | 0 |
| 173675 | 188389 | B009SF0TN6 | A1L0GWGRK4BYPT | Bety Robinson | 0 | 0 |
| 204727 | 221795 | B009SR4OQ2 | A32A6X5KCP7ARG | sicamar | 1 | 1 |
| 5259 | 5703 | B009WSNWC4 | AMP7K1O84DH1T | ESTY | 0 | 0 |
| 302474 | 327601 | B009WVB40S | A3ME78KVX31T21 | K'la | 0 | 0 |

364173 rows × 12 columns

```
final["Combined_Text"] = final["Cleaned_text"].values + ' ' + final['Cleaned_summary'].values
```

```
↳ /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/10min.html
"""Entry point for launching an IPython kernel.

final

```
↳
```


| | | | | | |
|--------|--------|------------|----------------|-----------------------------|-----|
| 138688 | 150506 | 0006641040 | A2IW4PEEKO2R0U | Tracy | 1 |
| 138689 | 150507 | 0006641040 | A1S4A3IQ2MU7V4 | sally sue "sally sue" | 1 |
| 138690 | 150508 | 0006641040 | AZGXZ2UUK6X | Catherine Hallberg "(Kate)" | 1 |
| 138691 | 150509 | 0006641040 | A3CMRKGE0P909G | Teresa | 3 |
| ... | ... | ... | ... | ... | ... |
| 178145 | 193174 | B009RSR8HO | A4P6AN2L435PV | romarc | 0 |
| 173675 | 188389 | B009SF0TN6 | A1L0GWGRK4BYPT | Bety Robinson | 0 |
| 204727 | 221795 | B009SR4OQ2 | A32A6X5KCP7ARG | sicamar | 1 |
| 5259 | 5703 | B009WSNWC4 | AMP7K1O84DH1T | ESTY | 0 |
| 302474 | 327601 | B009WVB40S | A3ME78KVX31T21 | K'la | 0 |

364173 rows × 13 columns

```
final.Combined_Text[138706]
```

```
↳ 'witty little book makes son laugh loud recite car driving along always sing refrain le
```

```
final.Cleaned_summary[138706]
```

```
↳ 'every book educational'
```

```
final.Cleaned_text[138706]
```

```
↳ 'witty little book makes son laugh loud recite car driving along always sing refrain le
```

```
time_sorted_data=final.sort_values("Time")
```

```
final_data=time_sorted_data.take(np.random.permutation(len(final)))[:100000]  
print(final_data.shape)
```

```
↳ (100000, 13)
```

```
final_data
```

```
↳
```

| | Id | ProductId | UserId | ProfileName | HelpfulnessNumerator | HelpfulnessDenominator |
|---------------|-----------|------------------|----------------|---------------------------------------------|-----------------------------|-------------------------------|
| 310032 | 335719 | B001VNGMMK | A3BZ51CB5PJ3VL | Jeffrey Smith | 1 | 1 |
| 33460 | 36400 | B004CLCEDE | A3S7PRF6YODY9N | luvallmykids | 1 | 1 |
| 278521 | 301805 | B0016BU7GO | A3KXIEF51W42PW | Cindy Luymes | 3 | 3 |
| 460187 | 497661 | B000HDK0D2 | ABHP4BWJBX9NT | Tim Ly | 5 | 5 |
| 515514 | 557343 | B0014C2JFC | A14S1X4IG0IJB | K. Strain "sMilesToGo" | 1 | 1 |
| ... | ... | ... | ... | ... | ... | ... |
| 377779 | 408499 | B004CQVCCS | A20BODDLOJMVQB | George T. Chambers III "George Chambers" | 0 | 0 |
| 171898 | 186486 | B001L4JH5I | AUV1H02P9GAAA | Starr Messer "Starr" | 0 | 0 |
| 356707 | 385834 | B000TV4W2C | A347PIFND2R6QW | K. Henderson | 1 | 1 |
| 477111 | 515946 | B000JSQKS4 | A27EMCN9BTYJHF | RoseMarie Cowham | 0 | 0 |
| 126220 | 136949 | B002AQP5FW | A2843500EKO5YB | Megan Gorg | 49 | 49 |

100000 rows × 13 columns

time_sorted_data



| | | | | | |
|---------------|--------|------------|----------------|----------------------------------|-----|
| 138706 | 150524 | 0006641040 | ACITT7DI6IDDL | shari zychinski | 0 |
| 138683 | 150501 | 0006641040 | AJ46FKXOVC7NR | Nicholas A Mesiano | 2 |
| 417839 | 451856 | B00004CXX9 | AIUWLEQ1ADEG5 | Elizabeth Medina | 0 |
| 346055 | 374359 | B00004CI84 | A344SMIA5JECGM | Vincent P. Ross | 1 |
| 417838 | 451855 | B00004CXX9 | AJH6LUC1UT1ON | The Phantom of the Opera | 0 |
| ... | ... | ... | ... | ... | ... |
| 511105 | 552637 | B0012EYELE | AWQOTHBNJBSVB | Gregory E. Grant "GG" | 0 |
| 282454 | 306008 | B0058CGLH6 | A3QRR5YN6ALFPG | james a riche | 0 |
| 311138 | 336872 | B0012KB4U2 | AGQBI6601XH2R | DaniC | 0 |
| 524273 | 566798 | B001PQTYN2 | A3OTHWG8LLCLMU | PACKERS FAN "Gordon Boone" | 1 |
| 355171 | 384161 | B000EVWQZW | A2PCNXBSKCABG5 | Whit | 0 |

364173 rows × 13 columns

X_features = final_data['Combined_Text'].values

```
y_target = final_data['Score']
```

```
x_train, x_test, y_train, y_test = train_test_split(X_features, y_target, test_size=0.25, r
```

```
↳
```

```
x_train.shape
```

```
↳ (75000,)
```

```
x_test.shape
```

```
↳ (25000,)
```

Featurization

```
from sklearn.feature_extraction.text import CountVectorizer
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
count_vec=CountVectorizer(min_df=10,max_features=50000,dtype="float")
```

```
count_vec.fit(x_train)
```

```
↳ CountVectorizer(analyzer='word', binary=False, decode_error='strict',  
dtype='float', encoding='utf-8', input='content',  
lowercase=True, max_df=1.0, max_features=50000, min_df=10,  
ngram_range=(1, 1), preprocessor=None, stop_words=None,  
strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',  
tokenizer=None, vocabulary=None)
```

```
bow_x_train=count_vec.transform(x_train)
```

```
bow_x_test=count_vec.transform(x_test)
```

```
print(bow_x_train.shape)
```

```
print(bow_x_test.shape)
```

```
↳ (75000, 11191)  
(25000, 11191)
```

```
bow_feature=count_vec.get_feature_names()
```

```
"""bow_x_train=bow_x_train.toarray()
```

```
bow_x_test=bow_x_test.toarray()"""
```

```
↳ 'bow_x_train=bow_x_train.toarray()\nbow_x_test=bow_x_test.toarray()'
```

```
bow_vocab=count_vec.vocabulary_
```

```
bow_vocab["not"]
```

```
#no of times words has been repeated in review corpus
```

```
↳ 6601
```

```

from sklearn.model_selection import TimeSeriesSplit
import xgboost
from xgboost import XGBClassifier

from sklearn.metrics import roc_auc_score,roc_curve,accuracy_score,confusion_matrix,classif

from sklearn.model_selection import RandomizedSearchCV

%%time
parameter={"n_estimators":[100,200,300,400,500],"max_depth":[3,4,5,6,7,8,9,10]}
cv_timeseries=TimeSeriesSplit(n_splits=5).split(bow_x_train)
xgb=XGBClassifier(n_jobs=-1)
xgb_bow=RandomizedSearchCV(xgb,param_distributions=parameter,scoring="roc_auc",cv=cv_timese
xgb_bow.fit(bow_x_train,y_train)

print("Best parameter obtained from RandomSearch CV: \n", xgb_bow.best_params_)
print("Best Score : ", xgb_bow.best_score_)

```

```

☐➤ Best parameter obtained from RandomSearch CV:
    {'n_estimators': 500, 'max_depth': 9}
    Best Score : 0.9558639370509707
    CPU times: user 1h 13min 27s, sys: 4.57 s, total: 1h 13min 32s
    Wall time: 19min 12s

```

```

def plot_model(title,clf,x_train,x_test,y_train,y_test):
    print(title)
    print("="*50)

    y_train_pred=clf.predict_proba(x_train)
    tr_fpr,tr_tpr,tr_thre=roc_curve(y_train,y_train_pred[:,1])
    tr_auc=auc(tr_fpr,tr_tpr)

    y_test_pred=clf.predict_proba(x_test)
    t_fpr,t_tpr,t_thres=roc_curve(y_test,y_test_pred[:,1])
    t_auc=auc(t_fpr,t_tpr)

    plt.plot([0,1],[0,1],"r--")
    plt.plot(tr_fpr,tr_tpr,label="tr_auc=%0.2f"%tr_auc)
    plt.plot(t_fpr,t_tpr,label="t_auc=%0.2f"%t_auc)
    plt.xlabel("fpr")
    plt.ylabel("tpr")
    plt.legend(loc="lower right")
    plt.title("ROC Curve")
    plt.show()


    y_pred=clf.predict(x_test)
    accuracy=accuracy_score(y_test,y_pred)
    confuse_mat=confusion_matrix(y_test,y_pred)
    cls_report=classification_report(y_test,y_pred)
    print("The Classification Report:\n",cls_report)

    sns.heatmap(confuse_mat,annot=True,fmt="g")
    plt.xlabel("Observed Value")
    plt.ylabel("Predicted Value")
    plt.title("Confusion Matrix")
    plt.show()

```


```
plot_model("XGB00ST WITH BAG OF WORDS",xgb_bow,bow_x_train,bow_x_test,y_train,y_test)
```

☞ XGB00ST WITH BAG OF WORDS



The Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.67 | 0.75 | 3842 |
| 1 | 0.94 | 0.98 | 0.96 | 21158 |
| accuracy | | | 0.93 | 25000 |
| macro avg | 0.90 | 0.82 | 0.86 | 25000 |
| weighted avg | 0.93 | 0.93 | 0.93 | 25000 |



```
!pip install chart_studio
```

☞



```
pip install wordcloud
```

```
↳ Requirement already satisfied: wordcloud in /usr/local/lib/python3.6/dist-packages (1.5.0)
Requirement already satisfied: pillow in /usr/local/lib/python3.6/dist-packages (from wordcloud) (6.2.0)
Requirement already satisfied: numpy>=1.6.1 in /usr/local/lib/python3.6/dist-packages (from wordcloud) (1.16.2)
```

```
from wordcloud import WordCloud
from IPython.core.display import HTML
```

```
bow_i_features=xgb_bow.best_estimator_.feature_importances_
```

```
↳ ERROR! Session/line number was not unique in database. History logging moved to new session.
```

```
bow_i_features
```

```
↳ array([0., 0., 0., ..., 0., 0., 0.], dtype=float32)
```

```
-bow_i_features
```

```
↳ array([-0., -0., -0., ..., -0., -0., -0.], dtype=float32)
```

```
top_features_index=(-bow_i_features).argsort()
```

```
(bow_i_features).argsort()
```

```
↳ array([    0,  7198,  7199, ..., 11060,  5795,  2606])
```

```
top_features_index
```

```
↳ array([ 2606,  5795, 11060, ...,  3995,  3974, 11190])
```

```
top_feature_bow=np.take(bow_feature,top_features_index[:20])
print(top_feature_bow)
```

```
↳ ['delicious' 'loves' 'worst' 'yummy' 'easy' 'horrible' 'perfect'
    'wonderful' 'excellent' 'awful' 'highly' 'great' 'yuck' 'terrible'
    'awesome' 'waste' 'beware' 'favorite' 'best' 'threw']
```

```
text=" "
```

```
for i in top_feature_bow:
    text=text+" "+i
```

```
text
```

```
↳ ' delicious loves worst yummy easy horrible perfect wonderful excellent awful highly g
```

```
def show_wordcloud(data, title = None):
```


```
    wordcloud = WordCloud(
        background_color='black',
        stopwords=stopwords,
        max_words=200,
        max_font_size=40,
        scale=3,
        random_state=1 # chosen at random by flipping a coin; it was heads
    ).generate(str(data))
```

```
fig = plt.figure(1, figsize=(12, 12))
plt.axis('off')
if title:
    fig.suptitle(title, fontsize=20)
    fig.subplots_adjust(top=2.3)
```

```
plt.imshow(wordcloud)
plt.show()
```

```
wc_bow=show_wordcloud(text,title="bow_top_features")
```

```
↳
```



bow_top_features

```
tf_vec=TfidfVectorizer(min_df=10,max_features=50000,dtype="float")
```

```
tf_vec.fit(x_train)
```

```
↳ /usr/local/lib/python3.6/dist-packages/sklearn/feature_extraction/text.py:1817: UserWarning
```

```
Only (<class 'numpy.float64'>, <class 'numpy.float32'>, <class 'numpy.float16'>) 'dtype'
```

```
TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype='float', encoding='utf-8', input='content',
lowercase=True, max_df=1.0, max_features=50000, min_df=10,
ngram_range=(1, 1), norm='l2', preprocessor=None,
smooth_idf=True, stop_words=None, strip_accents=None,
sublinear_tf=False, token_pattern='(?u)\\b\\w\\w+\\b',
tokenizer=None, use_idf=True, vocabulary=None)
```

```
tf_feature=tf_vec.get_feature_names()
```

```
tf_x_train=tf_vec.transform(x_train)
```

```
tf_x_test=tf_vec.transform(x_test)
```

```
print(tf_x_train.shape)
```

```
print(tf_x_test.shape)
```

```
↳ (75000, 11191)
(25000, 11191)
```

```
%%time
```

```
parameter={"n_estimators":[100,200,300,400,500],"max_depth":[3,4,5,6,7,8,9,10]}
```

```
cv_timeseries=TimeSeriesSplit(n_splits=5).split(tf_x_train)
```

```
xgb=XGBClassifier(n_jobs=-1)
```


```
xgb_tfidf=RandomizedSearchCV(xgb,param_distributions=parameter,scoring="roc_auc",cv=cv_time)
xgb_tfidf.fit(tf_x_train,y_train)
```

```
print("Best parameter obtained from RandomSearch CV: \n", xgb_tfidf.best_params_)
print("Best Score : ", xgb_tfidf.best_score_)
```

```
➤ Best parameter obtained from RandomSearch CV:
  {'n_estimators': 500, 'max_depth': 7}
Best Score : 0.9557485656814665
CPU times: user 2h 37min 33s, sys: 7.88 s, total: 2h 37min 41s
Wall time: 40min 17s
```


```
plot_model("XGB00ST with TFIDF",xgb_tfidf,tf_x_train,tf_x_test,y_train,y_test)
```

```
➤ XGB00ST with TFIDF
```



The Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.65 | 0.74 | 3842 |
| 1 | 0.94 | 0.98 | 0.96 | 21158 |
| accuracy | | | 0.93 | 25000 |
| macro avg | 0.90 | 0.81 | 0.85 | 25000 |
| weighted avg | 0.93 | 0.93 | 0.92 | 25000 |



```
def plot_3d(n_estimators,max_depth,tr_auc,te_auc):
    configure_plotly_browser_state()
    x1=n_estimators
    y1=max depth
```

```
z1=tr_auc
```

```
x2=n_estimators
```

```
y2=max_depth
```

```
z2=te_auc
```

```
offline.init_notebook_mode(connected=False)
```

```
trace1=go.Scatter3d(x=x1,y=y1,z=z1,name="train_auc")
```

```
trace2=go.Scatter3d(x=x2,y=y2,z=z2,name="cv_auc")
```

```
data=[trace1,trace2]
```

```
layout=go.Layout(scene=dict(xaxis = dict(title='n_estimators'),yaxis = dict(title='max_de
```

```
#layout=go.layout(scene=dict(xaxis=dict(title="n_estimators"),yaxis=dict(title="max_depth
```

```
fig = go.Figure(data=data, layout=layout)
```

```
#fig=go.Figure(data=data,layout=layout)
```

```
offline.iplot(fig, filename='3d-scatter-colorscale')
```

```
#offline.iplot(fig,filename="3d-scatter-colorspace")
```

```
feature=xgb_tfidf.best_estimator_.feature_importances_
```

```
top_features_index_tf=(-feature).argsort()
```

```
top_features_index_tf
```

```
↳ array([ 2606, 11060,  5795, ...,  3928,  3930, 11190])
```

```
top_feature_tfidf=np.take(tf_feature,top_features_index_tf[:100])
```

```
print(top_feature_tfidf)
```

```
↳ ['delicious' 'worst' 'loves' 'horrible' 'threw' 'awful' 'disappointed'
'yuck' 'beware' 'excellent' 'best' 'waste' 'misleading' 'favorite'
'perfect' 'wonderful' 'return' 'easy' 'yummy' 'terrible' 'disappointment'
'description' 'great' 'gross' 'disappointing' 'money' 'highly' 'tasty'
'yum' 'rip' 'love' 'refund' 'nice' 'trash' 'poor' 'smooth' 'awesome'
'stale' 'amazing' 'disgusting' 'glad' 'pleased' 'china' 'worse' 'happy'
'morning' 'expiration' 'overpriced' 'fantastic' 'works' 'sadly' 'add'
'ruined' 'refreshing' 'stores' 'weak' 'garbage' 'shame' 'bland' 'snack'
'label' 'bad' 'away' 'false' 'nasty' 'dissapointed' 'advertising'
'hooked' 'enjoy' 'thank' 'keeps' 'pleasantly' 'fda' 'tasteless' 'beat'
'always' 'easier' 'family' 'unfortunately' 'diarrhea' 'quickly' 'expired'
'ended' 'deceptive' 'wrong' 'artificial' 'mistake' 'watery' 'meal' 'paid'
'batch' 'barely' 'without' 'dangerous' 'chemical' 'inedible' 'opened'
'cold' 'definitely' 'stick']
```

```
text_tf=" "
```

```
for i in top_feature_tfidf:
```

```
    text_tf=text_tf+" "+i
```

```
wc_tf=show_wordcloud(text_tf,title="tfidf_top_features")
```

```
↳
```



```
return sent_vectors
```

```
j=["my name is sushil","my last name is chauhan"]  
i=0  
list_of_sentence=[]  
for sentence in j:  
    list_of_sentence.append(sentence.split())  
list_of_sentence
```

```
↳ [['my', 'name', 'is', 'sushil'], ['my', 'last', 'name', 'is', 'chauhan']]
```

```
%%time  
avg_x_train=avg_words(x_train)  
avg_x_test=avg_words(x_test)
```

```
↳ Buffered data was truncated after reaching the output size limit.
```

```
len(avg_x_train)
```

```
↳ 75000
```

```
a_x_train=np.array(avg_x_train)  
a_x_test=np.array(avg_x_test)
```

```
↳
```

```
a_x_train.shape
```

```
↳ (75000, 50)
```

```
a_x_test.shape
```

```
↳ (25000, 50)
```


```
%%time  
parameter={"n_estimators":[100,200,300,400,500],"max_depth":[3,4,5,6,7,8,9,10]}  
cv_timeseries=TimeSeriesSplit(n_splits=5).split(a_x_train)  
xgb=XGBClassifier(n_jobs=-1)  
xgb_avg=RandomizedSearchCV(xgb,param_distributions=parameter,scoring="roc_auc",cv=cv_timeseries)  
xgb_avg.fit(a_x_train,y_train)
```

```
print("Best parameter obtained from RandomSearch CV: \n", xgb_avg.best_params_)  
print("Best Score : ", xgb_avg.best_score_)
```

```
↳ Best parameter obtained from RandomSearch CV:  
    {'n_estimators': 400, 'max_depth': 8}  
Best Score : 0.9414421779936418  
CPU times: user 2h 14min 31s, sys: 6.06 s, total: 2h 14min 38s  
Wall time: 34min 10s
```


```
plot_model("XGB00ST WITH AVG_W2V",xgb_avg,a_x_train,a_x_test,y_train,y_test)
```

```
↳
```

The Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.80 | 0.19 | 0.30 | 3842 |
| 1 | 0.87 | 0.99 | 0.93 | 21158 |
| accuracy | | | 0.87 | 25000 |
| macro avg | 0.83 | 0.59 | 0.61 | 25000 |
| weighted avg | 0.86 | 0.87 | 0.83 | 25000 |



```
tr_auc=xgb_avg.cv_results_["mean_train_score"]
te_auc=xgb_avg.cv_results_["mean_test_score"]
plot_3d(n_estimators,max_depth,tr_auc,te_auc)
```



```

def tfidf(x_train):
    # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
    model = TfidfVectorizer()
    list_of_sentence=[]
    for sentence in x_train:
        list_of_sentence.append(sentence.split())

w2v_models=Word2Vec(list_of_sentence,size=50,min_count=10,workers=8)
w2v_words=list(w2v_models.wv.vocab)

tf_idf_matrix = model.fit_transform(x_train)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sentence: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_models.wv[word]

```

```

        #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
        # to reduce the computation we are
        # dictionary[word] = idf value of word in whole corpus
        # sent.count(word) = tf value of word in this review
        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1
return tfidf_sent_vectors

```

```

%%time
wtf_x_train=tfidf(x_train)
wtf_x_test=tfidf(x_test)

```

☞ Buffered data was truncated after reaching the output size limit.

```

w_x_train=np.array(wtf_x_train)
w_x_test=np.array(wtf_x_test)

```

☞

```

%%time
parameter={"n_estimators":[100,200,300,400,500],"max_depth":[3,4,5,6,7,8,9,10]}
cv_timeseries=TimeSeriesSplit(n_splits=5).split(w_x_train)
xgb=XGBClassifier(n_jobs=-1)
xgb_w=RandomizedSearchCV(xgb,param_distributions=parameter,scoring="roc_auc",cv=cv_timeseries)
xgb_w.fit(w_x_train,y_train)

```

```

print("Best parameter obtained from RandomSearch CV: \n", xgb_w.best_params_)
print("Best Score : ", xgb_w.best_score_)

```

☞ Best parameter obtained from RandomSearch CV:

```

{'n_estimators': 400, 'max_depth': 10}
Best Score : 0.9225804036534406
CPU times: user 2h 47min 25s, sys: 7.93 s, total: 2h 47min 33s
Wall time: 42min 20s


```

```

plot_model("XGBOOST with WEIGHTED TFIDF",xgb_w,w_x_train,w_x_test,y_train,y_test)

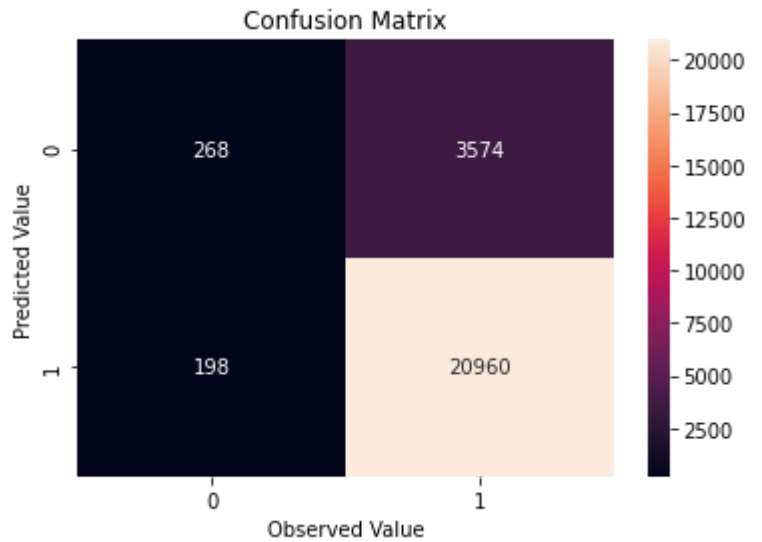
```

☞



The Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.58 | 0.07 | 0.12 | 3842 |
| 1 | 0.85 | 0.99 | 0.92 | 21158 |
| accuracy | | | 0.85 | 25000 |
| macro avg | 0.71 | 0.53 | 0.52 | 25000 |
| weighted avg | 0.81 | 0.85 | 0.80 | 25000 |




```
tr_auc=xgb_w.cv_results_["mean_train_score"]
ts_auc=xgb_w.cv_results_["mean_test_score"]
plot_3d(n_estimators,max_depth,tr_auc,ts_auc)
```



```
from prettytable import PrettyTable
```

```
ptable=PrettyTable()
```

```
ptable.add_column("XGB00ST",["BAG OF WORDS","TFIDF","AVERAGE-W2V","WEIGHTED-TFIDF"])
ptable.add_column("n_estimators",[500,500,400,400])
ptable.add_column("max_depth",[9,7,8,10])
ptable.add_column("train_auc",[0.99,0.99,1,1])
ptable.add_column("test_auc",[0.96,0.96,0.87,0.81])
ptable.add_column("Times in min",[19,40,34,42])
print(ptable)
```



| XGB00ST | n_estimators | max_depth | train_auc | test_auc | Times in min |
|----------------|--------------|-----------|-----------|----------|--------------|
| BAG OF WORDS | 500 | 9 | 0.99 | 0.96 | 19 |
| TFIDF | 500 | 7 | 0.99 | 0.96 | 40 |
| AVERAGE-W2V | 400 | 8 | 1 | 0.87 | 34 |
| WEIGHTED-TFIDF | 400 | 10 | 1 | 0.81 | 42 |

Conclusion:-