

```
In [1]: import nvforest
```

```
In [2]: lazy_imports()
```

```
Out[2]: ['import glob',
        'import datetime as dt',
        'import os',
        'import pickle',
        'import spacy',
        'from sklearn.model_selection import train_test_split',
        'import sys',
        'import numpy as np',
        'import matplotlib.pyplot as plt',
        'from sklearn.ensemble import RandomForestClassifier',
        'import nltk',
        'from sklearn.preprocessing import OneHotEncoder',
        'import statistics',
        'from dask import dataframe as dd',
        'import plotly as py',
        'import bokeh',
        'import plotly.graph_objs as go',
        'import gensim',
        'from sklearn.manifold import TSNE',
        'from sklearn.ensemble import RandomForestRegressor',
        'import seaborn as sns',
        'import plotly.express as px',
        'from sklearn.ensemble import GradientBoostingClassifier',
        'import pandas as pd',
        'from openpyxl import load_workbook',
        'import re',
        'from sklearn.feature_extraction.text import TfidfVectorizer',
        'import altair as alt',
        'from pyspark import SparkContext',
        'import dash',
        'import matplotlib as mpl',
        'from pathlib import Path',
        'import sklearn',
        'import tensorflow as tf',
        'from sklearn.ensemble import GradientBoostingRegressor',
        'import keras',
        'import tqdm',
        'import pydot',
        'from sklearn import svm']
```

```
In [3]: import sqlite3
```

```
In [4]: con=sqlite3.connect("/home/sushil/Downloads/Datasets/amazon-fine-food-reviews/database")
```

```
In [5]: database=con.cursor().execute("select * from reviews where Score<=3 limit 20000")
```

```
In [6]: database.execute("select Score, value, counts()")
```

```
Out[6]: 5      13745
        4       3110
        1       1953
        2       1192
        Name: Score, dtype: int64
```

```
In [7]: def partition(x):
        if x>3:
            return 1
        return 0
```

```
In [8]: database=database[database.HelpfulnessNumerator<=database.HelpfulnessDenominator]
```

```
In [9]: database
```

Out[9]:

	<b>Id</b>	<b>ProductId</b>	<b>UserId</b>	<b>ProfileName</b>	<b>HelpfulnessNumerator</b>	<b>HelpfulnessDenominator</b>	<b>Score</b>
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	1
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	3	3	1
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	1
...	...	...	...	...	...	...	...
19995	21784	B000KV61FC	A3FVKI0UH9DO2A	S. Malosh	1	1	1
19996	21785	B000KV61FC	A3ACVJEAM4L2LQ	Elb	1	1	1
19997	21786	B000KV61FC	AHHWZ4723VGOL	Sherry Lynn	1	1	1
19998	21787	B000KV61FC	A2O4CZ102I8Q2K	jus42day	1	1	1
19999	21788	B000KV61FC	A3I4GCI6XTX1BB	Eric C. Vizinas "Q"	1	1	1

20000 rows × 10 columns

```
In [10]: database.Score=database.Score.map(partition)
```

```
In [11]: database.Score.value_counts()
```

```
Out[11]: 1    16855
0       3145
Name: Score, dtype: int64
```

Datasets is Severely imbalance

```
In [12]: df=database.sort_values("ProductId")
```

```
In [13]: df=df.drop_duplicates(["ProductId","UserId","Time","Text"])
```

```
In [15]: df.shape
```

```
Out[15]: (19963, 10)
```

```
In [16]: tvne(df)
```

```
Out[16]: pandas.core.frame.DataFrame
```

#### Process Text Summary and Text Data Cleaning

```
In [17]: def preprocess_text(sentence):
def remove_html(sentence):
    html_tag_re_obj = re.compile('<.*>?')
    return re.sub(html_tag_re_obj, ' ', sentence)

def remove_punctuations(sentence):
    cleaned_sentence = re.sub(r'[^a-zA-Z]', r' ', sentence)
    return cleaned_sentence

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase

from bs4 import BeautifulSoup
import re
#from nltk.corpus import stopwords
import nltk
from tqdm import tqdm
from nltk.stem import SnowballStemmer
from nltk.corpus import stopwords
stopwords = stopwords.words('english')
stemmer = SnowballStemmer('english')
stopwords = set(stopwords)
stopwords.remove('not')

cleaned_corpus = []
for doc in sentence:
    cleaned_doc_1 = remove_html(doc)
    cleaned_doc_2 = remove_punctuations(doc)
    cleaned_doc_2 = decontracted(cleaned_doc_2)
    cleaned_corpus.append(cleaned_doc_2)
count = 0

filtered_corpus = list(map(lambda doc: ' '.join(list(filter(lambda word: True if word not in stopwords, doc.split()))), cleaned_corpus))
process_text = list(map(lambda doc: ' '.join(list(map(stemmer.stem, doc.split()))), filtered_corpus))

return process_text
```

```
In [18]: %%time
df.Text=preprocess_text(df.Text)
df.Summary=preprocess_text(df.Summary)

CPU times: user 18.6 s, sys: 252 ms, total: 18.9 s
Wall time: 25.1 s
```

```
In [19]: df["Text"]
```

```
Out[19]: 2547    we use victor fli bait season can beat great p...
2546    whi product avail br http www amazon com victo...
1146    this realli good idea final product outstand i...
1145    i receiv shipment could hard wait tri product ...
8696    i use brand year if feel clog ate massiv meal ...

...

11590    some product work need done disc i wast sever ...
11588    i love coffe usual i would not recommend maxwe...
11589    this cappuccino good tast almost good starbuck...
1362    this coffe suppos premium tast wateri thin not...
5259    purchas product local store ny kid love it qui...
Name: Text, Length: 19963, dtype: object
```

```
In [20]: df.Summary
```

```
Out[20]: 2547    fli begon
2546    thirti buck
1146    great product
1145    wow make slicker
8696    the best cleans tea i ever had

...

11590    cappucino disc
11588    coff shop qualiti capuccino home
11589    veri veri good
1362    weak coffe not good premium product price
5259    delici
Name: Summary, Length: 19963, dtype: object
```

```
In [21]: %%time
text_length=[]
for i in df.Text:
    text_length.append(len(i.split()))

CPU times: user 91.4 ms, sys: 3.88 ms, total: 95.3 ms
Wall time: 94.1 ms
```

```
In [22]: len(text_length)
```

```
Out[22]: 19963
```

```
In [23]: text_length[0]
```

```
Out[23]: 10
```

```
In [25]: type(df.Text)
```

```
Out[25]: pandas.core.series.Series
```

```
In [26]: df.shape
```

```
Out[26]: (19963, 10)
```

```
In [27]: index=list(df.index)
print(type(index))
<class 'list'>
```

```
In [28]: df.text_length=pd.Series(data=text_length,name="text_length",index=index)
```

```
In [29]: df.text_length
```

```
Out[29]: 2547      10
          2546      32
          1146      19
          1145      47
          8696      40
          ...
          11590     25
          11588     116
          11589      79
          1362      35
          5259      22
          Name: text_length, Length: 19963, dtype: int64
```

```
In [30]: df_new=pd.concat([df,df_text_length],axis=1)
df_new
```

```
Out[30]:
```

	<b>Id</b>	<b>ProductId</b>	<b>UserId</b>	<b>ProfileName</b>	<b>HelpfulnessNumerator</b>	<b>HelpfulnessDenominator</b>	<b>Sco</b>
<b>2547</b>	2775	B00002NCJC	A13RRPGE79XFFH	reader48	0	0	
<b>2546</b>	2774	B00002NCJC	A196AJHU9EASJN	Alex Chaffee	0	0	
<b>1146</b>	1245	B00002Z754	A29Z5PI9BW2PU3	Robbie	7	7	
<b>1145</b>	1244	B00002Z754	A3B8RCEI0FXFI6	B G Chase	10	10	
<b>8696</b>	9527	B00005V3DC	A8KY7S48EW7LW	A. Daly "AD"	0	0	
...	...	...	...	...	...	...	...
<b>11590</b>	12621	B009KP6HBM	A3UCO959VA9MV	Maryann Wardach	0	0	
<b>11588</b>	12619	B009KP6HBM	A20NB4UBW4WDKG	Gerardo "GD"	1	1	
<b>11589</b>	12620	B009KP6HBM	A3D9NUCR4RXDPY	Kathleen San Martino	1	1	
<b>1362</b>	1478	B009UOFU20	AJVB004EB0MVK	D. Christofferson	0	0	
<b>5259</b>	5703	B009WSNWC4	AMP7K1O84DH1T	ESTY	0	0	

19963 rows × 11 columns

```
In [31]: a=df_new.text_length.isnull()
```

```

In [32]: a.value_counts()
Out[32]: False      19963
         Name: text_length, dtype: int64

In [33]: df_new=df_new.sort_values("Time")

In [34]: x_train,x_test,y_train,y_test=train_test_split(df_new.Text,df_new.Score,test_size=0.33)

In [35]: print(x_train.shape)
         print(y_train.shape)
         (13375,)
         (13375,)

In [36]: from sklearn.feature_extraction.text import CountVectorizer
         from sklearn.feature_extraction.text import TfidfVectorizer

In [37]: from sklearn.model_selection import GridSearchCV

In [38]: count_vec=CountVectorizer(min_df=10)

In [39]: count_vec.fit(x_train)
Out[39]: CountVectorizer(analyzer='word', binary=False, decode_error='strict',
                        dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
                        lowercase=True, max_df=1.0, max_features=None, min_df=10,
                        ngram_range=(1, 1), preprocessor=None, stop_words=None,
                        strip_accents=None, token_pattern='(?u)\\b\\w\\w+\\b',
                        tokenizer=None, vocabulary=None)

In [40]: bow_x_train=count_vec.transform(x_train)
         bow_x_test=count_vec.transform(x_test)

In [43]: bow_features=count_vec.get_feature_names()
         bow_vocab=count_vec.vocabulary_

In [41]: print(bow_x_train.shape)
         print(bow_x_test.shape)
         (13375, 3424)
         (6588, 3424)

In [42]: bow_x_train=bow_x_train.toarray()
         bow_x_test=bow_x_test.toarray()

In [52]: rm=RandomForestClassifier(n_jobs=-1,class_weight="balanced")

In [53]: parameter={"n_estimators":[100,200,300,400,500]}

In [54]: clf_bow=GridSearchCV(estimator=rm,param_grid=parameter,cv=4,scoring="roc_auc",return_t

In [55]: from sklearn.metrics import roc_auc_score,roc_curve,classification_report,confusion_m

```

In [56]: %%time

```
clf_bow_fit(bow_x_train, y_train)
```

CPU times: user 10min 42s, sys: 11.1 s, total: 10min 53s

Wall time: 24min 7s

```
Out[56]: GridSearchCV(cv=4, error_score=nan,
                      estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
                                                         class_weight='balanced',
                                                         criterion='gini', max_depth=None,
                                                         max_features='auto',
                                                         max_leaf_nodes=None,
                                                         max_samples=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=2,
                                                         min_weight_fraction_leaf=0.0,
                                                         n_estimators=100, n_jobs=-1,
                                                         oob_score=False,
                                                         random_state=None, verbose=0,
                                                         warm_start=False),
                      iid='deprecated', n_jobs=None,
                      param_grid={'n_estimators': [100, 200, 300, 400, 500]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring='roc_auc', verbose=0)
```

In [57]: print(clf\_bow.best\_params\_)

```
{'n_estimators': 500}
```

In [58]: print(clf\_bow.best\_score\_)

```
0.9221189236360187
```



```

In [69]: def plot_model(title,clf,x_train,y_train,x_test,y_test):

    print(title)
    print("="*100)

    y_train_pred=clf.predict_proba(x_train)
    train_fpr,train_tpr,train_thres=roc_curve(y_train,y_train_pred[:,1])
    train_auc_score=auc(train_fpr,train_tpr)
    print("train_auc_score:",train_auc_score)
    y_test_pred=clf.predict_proba(x_test)
    test_fpr,test_tpr,test_thres=roc_curve(y_test,y_test_pred[:,1])
    test_auc_score=auc(test_fpr,test_tpr)
    print("test_auc_score",test_auc_score)
    plt.plot(train_fpr,train_tpr,label="train auc =%0.2f"+str(train_auc_score))
    plt.plot(test_fpr,test_tpr,label="test auc = %0.2f"+str(test_auc_score))
    plt.plot([0,1],[0,1],"r--")
    plt.xlabel("fpr")
    plt.ylabel("tpr")
    plt.legend("bottom right")
    plt.title("ERROR PLOTS")
    plt.show()

    y_pred_test=clf.predict(x_test)
    accuracy=accuracy_score(y_test,y_pred_test)
    print("the accuracy score of our model on test:",accuracy)

    class_report=classification_report(y_test,y_pred_test)
    print("the classification reports \n",class_report)

    confuse_mat=confusion_matrix(y_test,y_pred_test)
    print("the confusion matrix :\n",confuse_mat)

    sns.heatmap(confuse_mat,annot=True,fmt="g")
    plt.xlabel("observed value")
    plt.ylabel("predict value")
    plt.show()

```

```
In [70]: plot_model("Random Forest on Bag of words" clf bow bow x_train y_train bow x_test y_test)
```


Random Forest on Bag of words

=====

=====

train\_auc\_score: 1.0

test\_auc\_score 0.9345570555570866




the accuracy score of our model on test: 0.8838797814207651

the classification reports

	precision	recall	f1-score	support
0	0.89	0.31	0.46	1048
1	0.88	0.99	0.93	5540
accuracy			0.88	6588
macro avg	0.89	0.65	0.70	6588
weighted avg	0.88	0.88	0.86	6588

the confusion matrix :

```
[[ 323  725]
 [   40 5500]]
```



```
In [71]: tf_vec=TfidfVectorizer(min_df=10)
```

```
In [72]: tf_vec.fit(x_train)
```

```
Out[72]: TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',  
                        dtype=<class 'numpy.float64'>, encoding='utf-8',  
                        input='content', lowercase=True, max_df=1.0, max_features=None,  
                        min_df=10, ngram_range=(1, 1), norm='l2', preprocessor=None,  
                        smooth_idf=True, stop_words=None, strip_accents=None,  
                        sublinear_tf=False, token_pattern='(?u)\\b\\w+\\b',  
                        tokenizer=None, use_idf=True, vocabulary=None)
```

```
In [73]: tfidf_x_train=tf_vec.transform(x_train).toarray()  
         tfidf_x_test=tf_vec.transform(x_test).toarray()
```

```
In [74]: tfidf_features_names=tf_vec.get_feature_names()
```


```
In [75]: tfidf_vocab=tf_vec.vocabulary
```

```
In [76]: %%time  
         clf_tfidf=GridSearchCV(estimator=rm,param_grid=parameter,cv=4,scoring="roc_auc",return  
         clf_tfidf.fit(tfidf_x_train,y_train)  
         print(clf_tfidf.best_params_)  
         print(clf_tfidf.best_score_)  
  
         {'n_estimators': 300}  
         0.9241887548388995  
         CPU times: user 7min 45s, sys: 14.3 s, total: 7min 59s  
         Wall time: 26min 36s
```

```
In [77]: plot_model("Random forest with tfidf" clf tfidf tfidf x_train y_train tfidf x_test y_test)
```

Random forest with tfidf

```
=====
train_auc_score: 1.0
test_auc_score 0.9327563073552512
```




the accuracy score of our model on test: 0.8816029143897997

the classification reports

	precision	recall	f1-score	support
0	0.90	0.29	0.44	1048
1	0.88	0.99	0.93	5540
accuracy			0.88	6588
macro avg	0.89	0.64	0.68	6588
weighted avg	0.88	0.88	0.85	6588

the confusion matrix :

```
[[ 301  747]
 [   33 5507]]
```



```
In [79]: from gensim.models import Word2Vec
```

```
In [146]: list_of_sentence_train=[]
for sentence in x_train:
    list_of_sentence_train.append(sentence.split())
```

```
In [147]: print(list(sentence_train[1]))
```

```
['glutino', 'gluten', 'free', 'cracker', 'proabl', 'closest', 'thing', 'get', 'ritz',  
'club', 'like', 'cracker', 'these', 'cracker', 'crisp', 'without', 'feel', 'stale',  
'nice', 'crunch', 'i', 'miss', 'regular', 'cracker', 'the', 'flavor', 'decent', 'i',  
'never', 'ate', 'cracker', 'alon', 'i', 'still', 'i', 'judg', 'tast', 'well', 'p',  
air', 'dip', 'spread', 'chees', 'the', 'flavor', 'subtl', 'enough', 'put', 'virtual',  
'anyth', 'top', 'be', 'warn', 'get', 'stale', 'quick', 'keep', 'zip', 'lock', 'bag',  
'even', 'box', 'the', 'price', 'cours', 'anoth', 'downsid', 'like', 'gluten', 'free',  
ree', 'food', 'i', 'develop', 'avers', 'price', 'gluten', 'free', 'product']
```

```
In [158]: w2v_models=Word2Vec(list(sentence_train) size=100 min_count=10 workers=4)
```

```
In [159]: w2v_models.vocabulary
```

```
Out[159]: <gensim.models.word2vec.Word2VecVocab at 0x7f3c581a07d0>
```

```
In [160]: w2v_words=list(w2v_models.wv.vocab)  
print(w2v_words[:50])
```

```
['robust', 'aromat', 'time', 'pop', 'contain', 'top', 'last', 'drop', 'pot', 'br', 'great',  
'flavor', 'black', 'mix', 'cream', 'sugar', 'glutino', 'gluten', 'free', 'cracker',  
'proabl', 'closest', 'thing', 'get', 'ritz', 'club', 'like', 'these', 'crisp',  
'without', 'feel', 'stale', 'nice', 'crunch', 'i', 'miss', 'regular', 'the', 'decent',  
'never', 'ate', 'alon', 'still', 'judg', 'tast', 'well', 'pair', 'dip', 'spread',  
'chees']
```

```
In [151]: w2v_models.wv["robust"]
```

```
Out[151]: array([ 6.51459172e-02, -7.78796613e-01,  2.80000001e-01, -1.15873307e-01,  
                  2.27104917e-01, -2.63447672e-01,  5.04275877e-03, -3.60053539e-01,  
                 -1.91085145e-01, -7.23941207e-01,  1.03078105e-01, -2.45274186e-01,  
                 -2.60765910e-01, -8.56026649e-01,  2.03870103e-01, -2.66229630e-01,  
                  3.71845365e-01,  6.93805158e-01,  2.10332870e-02,  5.75901031e-01,  
                  1.00094378e+00,  1.15692258e-01, -5.58268130e-01, -7.22508907e-01,  
                 -8.95051509e-02, -4.06676948e-01,  5.56860030e-01, -2.79707164e-02,  
                 -5.93563676e-01,  3.42972249e-01, -3.46264452e-01, -2.49021590e-01,  
                 -9.85664546e-01, -6.12769246e-01,  5.50739467e-01, -5.50841451e-01,  
                 -5.80104709e-01,  1.66340202e-01, -3.29534203e-01,  8.75152588e-01,  
                 -5.43673038e-01, -1.09755814e+00, -6.01606965e-01,  3.54879647e-01,  
                 -9.77566764e-02, -1.35977857e-03, -3.69163692e-01, -2.13634953e-01,  
                  1.79683864e-01, -3.60152453e-01,  4.48719949e-01, -1.09393764e+00,  
                 -2.32582822e-01, -3.18637043e-01,  7.39820957e-01, -8.12138379e-01,  
                 -4.40511376e-01,  1.78786069e-01,  4.65160161e-01,  5.63497901e-01,  
                 -1.10800803e-01,  1.33539056e-02,  5.83385348e-01,  8.72598946e-01,  
                 -9.72747803e-01, -4.75852460e-01,  2.35692393e-02, -9.43671286e-01,  
                 -4.77544785e-01, -4.81497735e-01,  2.42418051e-01,  2.17397541e-01,  
                 -3.04513387e-02,  4.01639372e-01, -8.15329850e-01,  4.81980950e-01,  
                 -3.94194067e-01, -1.30939376e+00,  4.24000621e-02,  4.90813315e-01,  
                 -1.35115814e+00,  1.40824854e+00, -3.54114711e-01,  5.31737983e-01,  
                 -3.55638117e-01, -2.38751635e-01, -6.91885173e-01,  1.78309411e-01,  
                  1.71968591e+00, -4.46066111e-01,  5.94441712e-01, -5.32339036e-01,  
                 -1.45108953e-01,  2.74965495e-01,  8.53092849e-01,  7.79579818e-01,  
                  2.42316425e-01,  1.04311742e-02,  9.84088480e-01,  3.21566433e-01],  
                dtype=float32)
```

```
In [161]: def avg_words2vec(list_of_sentence,w2v_model,w2v_words):
    sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
    for sent in list_of_sentence: # for each review/sentence
        sent_vec = np.zeros(100) # as word vectors are of zero length 50, you might not
        cnt_words = 0; # num of words with a valid vector in the sentence/review
        for word in sent: # for each word in a review/sentence
            if word in w2v_words:
                vec = w2v_model.wv[word]
                sent_vec += vec
                cnt_words += 1
        if cnt_words != 0:
            sent_vec /= cnt_words
        sent_vectors.append(sent_vec)
    sent_vectors_train = np.array(sent_vectors)
    return sent_vectors,time
```

```
In [153]: list_of_sentence_test=[]
    for sentence in x_test:
        list_of_sentence_test.append(sentence.split())
```

```
In [154]: w2v_model_test=Word2Vec(list_of_sentence_test,size=100,min_count=10,workers=4)
```

```
In [155]: w2v_words_test=list(w2v_model_test.wv.vocab)
```

```
In [169]: def avg_w2v(x_train):
    i=0
    list_of_sentence=[]
    for sentence in x_train:
        list_of_sentence.append(sentence.split())

    w2v_model=Word2Vec(list_of_sentence,size=100,min_count=10,workers=4)
    w2v_words=list(w2v_model.wv.vocab)
    # average Word2Vec
    # compute average word2vec for each review.
    sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
    for sent in list_of_sentence: # for each review/sentence
        sent_vec = np.zeros(100) # as word vectors are of zero length 50, you might not
        cnt_words = 0; # num of words with a valid vector in the sentence/review
        for word in sent: # for each word in a review/sentence
            if word in w2v_words:
                vec = w2v_model.wv[word]
                sent_vec += vec
                cnt_words += 1
        if cnt_words != 0:
            sent_vec /= cnt_words
        sent_vectors.append(sent_vec)
    sent_vectors = np.array(sent_vectors)
    return sent_vectors
```

```
In [170]: avg_x_train=avg_w2v(x_train)
    avg_x_test=avg_w2v(x_test)
```

```
In [171]: avg_x_train.shape
```

```
Out[171]: (13375, 100)
```

```
In [173]: avg_x_test.shape
```

```
Out[173]: (6588, 100)
```


```
In [174]: %%time
clf_avg=GridSearchCV(estimator=rm,param_grid=parameter,cv=4,scoring="roc_auc",return_t
clf_avg.fit(avg_x_train,y_train)
print(clf_tfidf.best_params_)
print(clf_tfidf.best_score_)

{'n_estimators': 300}
0.9241887548388995
CPU times: user 2min 20s, sys: 2.86 s, total: 2min 23s
Wall time: 4min 55s
```

```
In [181]: plot_model("RandomForest with Avg-words2vec" clf_avg_avg, x_train, y_train, avg, x_test, y_test)
```

RandomForest with Avg-words2vec

```
=====
train_auc_score: 1.0
test_auc_score 0.6421158059360101
```



the accuracy score of our model on test: 0.840922890103218

the classification reports


	precision	recall	f1-score	support
0	0.00	0.00	0.00	1048
1	0.84	1.00	0.91	5540
accuracy			0.84	6588
macro avg	0.42	0.50	0.46	6588
weighted avg	0.71	0.84	0.77	6588

the confusion matrix :

```
[[ 0 1048]
 [ 0 5540]]
```

/home/sushil/anaconda3/lib/python3.7/site-packages/sklearn/metrics/\_classification.py:1272: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```





```

In [187]: def weighted_tfidf(x_train):
            model = TfidfVectorizer()
            model.fit(x_train)
            i=0
            list_of_sentence=[]
            for sentence in x_train:
                list_of_sentence.append(sentence.split())

            w2v_model=Word2Vec(list_of_sentence,size=100,min_count=10,workers=4)
            w2v_words=list(w2v_model.wv.vocab)
            # we are converting a dictionary with word as a key, and the idf as a value
            dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
            tfidf_feat = model.get_feature_names() # tfidf words/col-names
            # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = t

            tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in the
            row=0;
            for sent in list_of_sentence: # for each review/sentence
                sent_vec = np.zeros(100) # as word vectors are of zero length
                weight_sum =0; # num of words with a valid vector in the sentence/review
                for word in sent: # for each word in a review/sentence
                    if word in w2v_words and word in tfidf_feat:
                        vec = w2v_model.wv[word]
                        #
                        tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
                        # to reduce the computation we are
                        # dictionary[word] = idf value of word in whole corpus
                        # sent.count(word) = tf value of word in this review
                        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
                        sent_vec += (vec * tf_idf)
                        weight_sum += tf_idf
                if weight_sum != 0:
                    sent_vec /= weight_sum
                tfidf_sent_vectors.append(sent_vec)
                row += 1

            return tfidf_sent_vectors

```

```

In [198]: w_x_train=weighted_tfidf(x_train)
           w_x_test=weighted_tfidf(x_test)

```

```

In [199]: w_x_train=np.array(w_x_train)

```

```

In [200]: w_x_test=np.array(w_x_test)

```

```

In [201]: print(w_x_test.shape)
           print(w_x_train.shape)
(6588, 100)
(13375, 100)

```

```


In [189]: %%time
           clf_w=GridSearchCV(estimator=rm,param_grid=parameter,cv=4,scoring="roc_auc",return_train_score=True)
           clf_w.fit(w_x_train,y_train)
           print(clf_w.best_params_)
           print(clf_w.best_score_)
{'n_estimators': 300}
0.9241887548388995
CPU times: user 1min 58s, sys: 2.55 s, total: 2min 1s
Wall time: 4min 56s

```

```
In [202]: plot_model("RandomForest with WeightedTFIDF" clf_w w x train v train w x test v test)
```

RandomForest with WeightedTFIDF

```
=====
train_auc_score: 0.8501717179667327
test_auc_score 0.6744074151900129
```



the accuracy score of our model on test: 0.840922890103218

the classification reports


	precision	recall	f1-score	support
0	0.00	0.00	0.00	1048
1	0.84	1.00	0.91	5540
accuracy			0.84	6588
macro avg	0.42	0.50	0.46	6588
weighted avg	0.71	0.84	0.77	6588

/home/sushil/anaconda3/lib/python3.7/site-packages/sklearn/metrics/\_classification.py:1272: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

the confusion matrix :

```
[[ 0 1048]
 [ 0 5540]]
```



```
In [205]: from prettytable import PrettyTable
```

```
In [217]: ntable1=PrettyTable()
```

```
In [218]: train_auc_score=[1,1,1,0.85]
test_auc_score=[0.93,0.92,0.64,0.67]
n_neighbors=[500,300,300,300]
names=["bag_of_words" "tfidf" "Avg-w2v" "weighted-tfidf"]
```

```
In [219]: ptable1.add_column("Random-Forest",names)
ptable1.add_column("n_neighbors",n_neighbors)
ptable1.add_column("TRAIN AUC",train_auc_score)
ptable1.add_column("TEST AUC",test_auc_score)
print(ntable1)
```

Random-Forest	n_neighbors	TRAIN AUC	TEST AUC
bag_of_words	500	1	0.93
tfidf	300	1	0.92
Avg-w2v	300	1	0.64
weighted-tfidf	300	0.85	0.67

```
In [ ]:
```